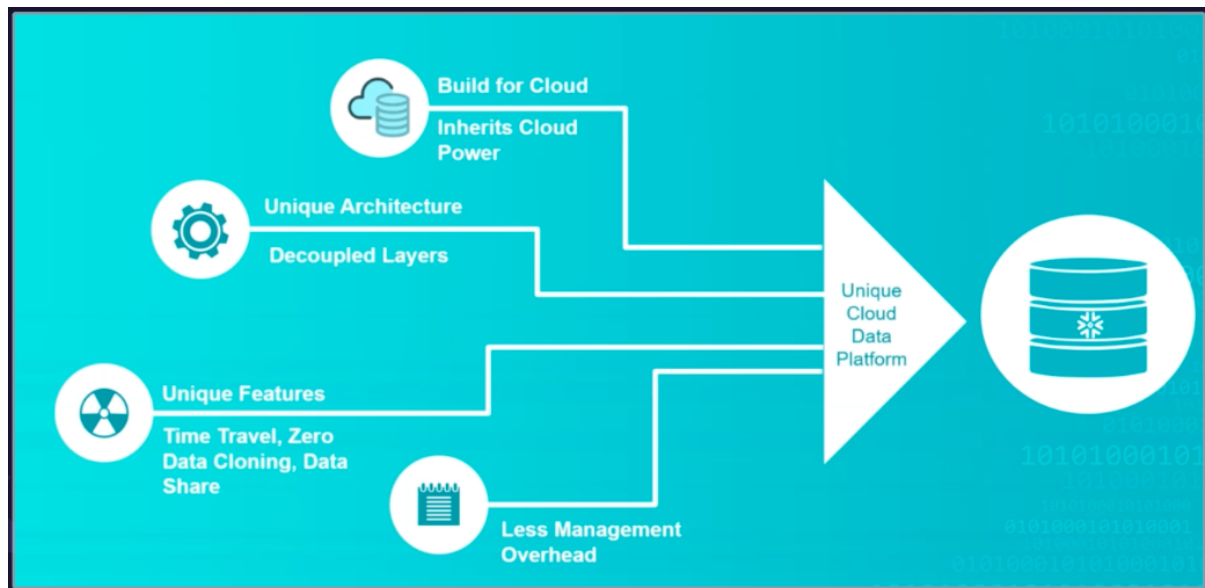


SNOWFLAKES

- Snowflake is an analytic data warehouse provided as Software-as-Services(SaaS). Snowflake provides data warehouse that is faster, easier to use and more flexible than other traditional data warehouses.
- Snowflake data warehouse is not built on existing databases or not on big data software platform as Hadoop.
- The snowflake data warehouse uses a new SQL database engine with unique architecture designed for the cloud.

Snowflakes key differentiator

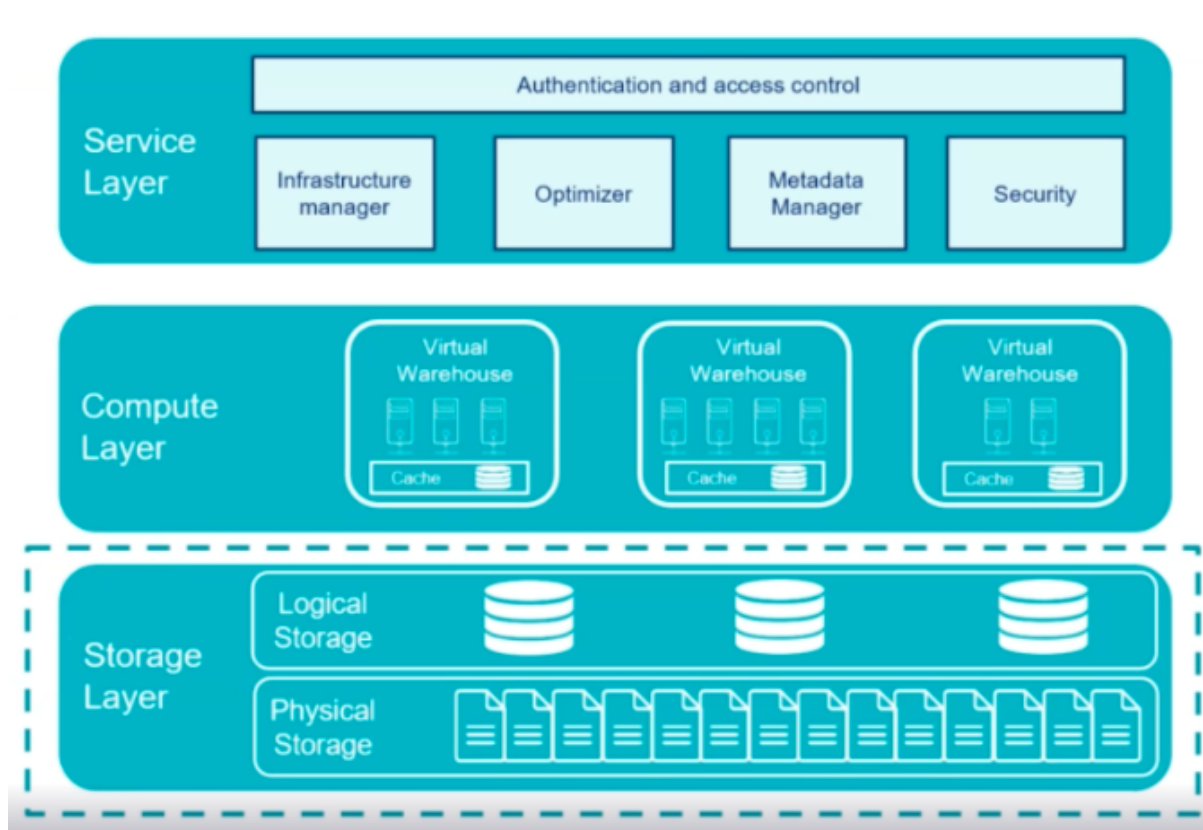


Key Concept and Architecture

Data Warehouse as Cloud Service:

- Snowflake data warehouse is true SaaS offering :
 - There is no hardware (virtual or physical) for you to select, install, configure and manage.
 - There is no software for you install, configure and manage.
 - Ongoing maintenance, management and tuning is handled by snowflake
- Snowflake completely runs on cloud infrastructure. All the component of the snowflake service runs on public cloud infrastructure
- Snowflake uses virtual compute instance for its compute need and storage service for storage of data. Snowflake can not be run on private cloud infrastructure(on premises)

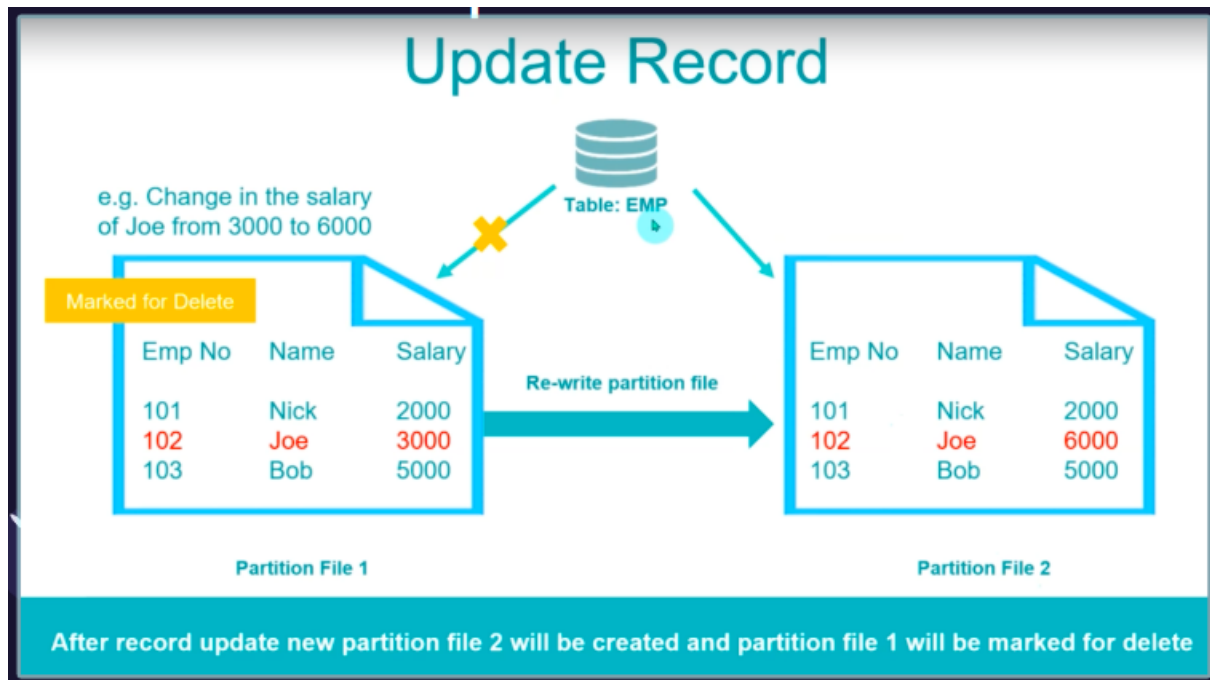
SNOWFLAKE ARCHITECTURE



Storage Layer

- This layer is independent from other 2 layer.
- Virtually unlimited storage capacity because it is connected to different cloud service provider eg. AWS, azure etc. in the backend.
- Data physically stored as micro partition files.
- Each micro partitioned file is of size 16 mb.
- Every file gets replicated 3 times to ensure high availability.
- File format is propriety to snowflake means snowflake has not exposed in which format it save file.
- File are immutable
- Table definition are logical and stored in its metadata layer. Here logical means it is there but u can't see it eg. underground metro u know there it is , but can't see it but u can feel it.

UPDATE Record



The mark for delete is done by snowflake itself and how it does it, that is not disclosed to the real world.

Compute (a.k.a Warehouse) Layer

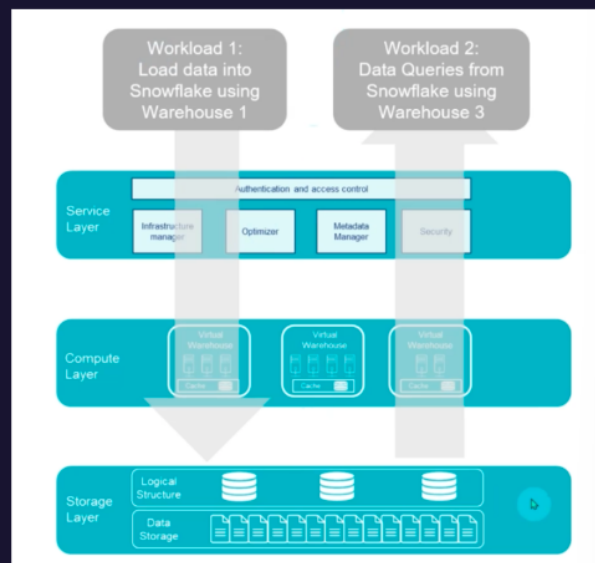
- Only layer to access data from storage layer
- It is scalable
 - Change the size of the warehouse at runtime
 - Add more nodes at the runtime
- Warehouse size varies from X-small (1 server/cluster) to 4X-Large (128 server/cluster)
- Warehouses are independent of each other.
- Every warehouse has its own small storage, which it uses to cache query data for better performance
- Compute can be auto-suspended when idle.

Warehouse Size & Pricing

Warehouse Size	Server/Cluster	Credit/Hour
X-Small	1	1
Small	2	2
Medium	4	4
Large	8	8
X-Large	16	16
2X-Large	32	32
3X-Large	64	64
4X-Large	128	128

	Standard	Enterprise	Business Critical
Cost per credit	\$2.20	\$3.30	\$4.40

Separate Warehouse for Different Workloads



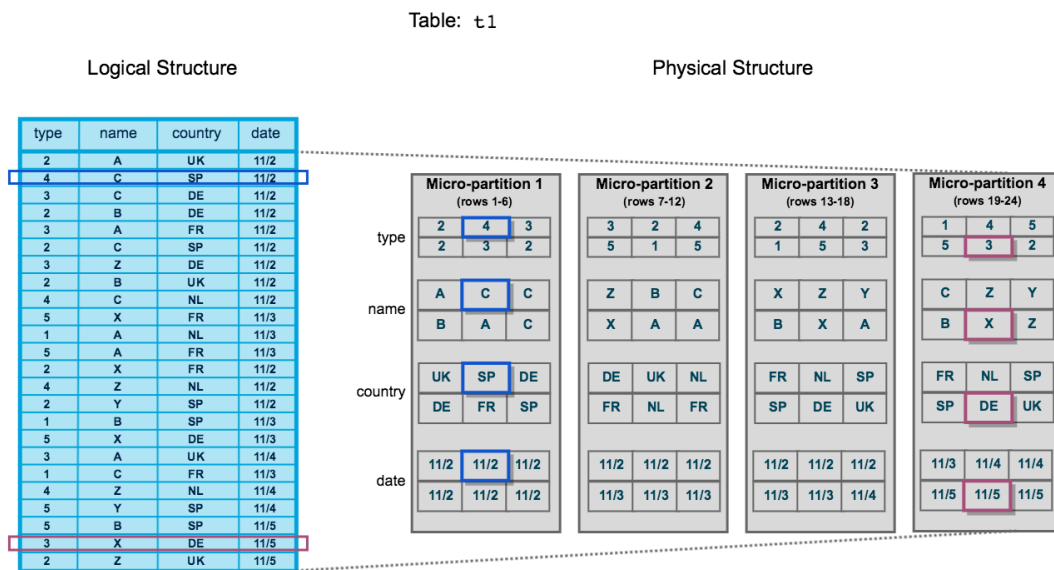
Service Layer

- Brain of the Snowflake.
- All the interaction with snowflake established via this layer
- Stores query output in a result cache
- Snowflake doesn't expose its service layer and users can't get insight and access their metadata.

- During data load service layer keeps track of which data stored in which partition file
- Maintain transaction consistency across Warehouse.

What are Micro partitions?

- All data in Snowflake tables is automatically divided into micro partitions, which are contiguous units of storage.
- Each Micro partitions contains between 50 mb to 500 mb of uncompressed data (note that the actual size in snowflake is smaller because data is always stored compressed.)
- Group of rows in tables are mapped into individual micro partitions, organized in a columnar fashion.



- Snowflake stores metadata about all rows stored in a micro-partition, including:
 - The range of values for each of the columns in the micro-partition.
 - The number of distinct values.
 - Additional properties used for both optimization and efficient query processing.

Note

Micro-partitioning is automatically performed on all Snowflake tables.

Tables are transparently partitioned using the ordering of the data as it is inserted/loaded.

Benefits of Micro-partitioning

The benefits of Snowflake's approach to partitioning table data include:

- In contrast to traditional static partitioning, Snowflake micro-partitions are derived automatically; they don't need to be explicitly defined up-front or maintained by users.
- As the name suggests, micro-partitions are small in size (50 to 500 MB, before compression), which enables extremely efficient DML and fine-grained pruning for faster queries.
- Micro-partitions can overlap in their range of values, which, combined with their uniformly small size, helps prevent skew.
- Columns are stored independently within micro-partitions, often referred to as *columnar storage*. This enables efficient scanning of individual columns; only the columns referenced by a query are scanned.
- Columns are also compressed individually within micro-partitions. Snowflake automatically determines the most efficient compression algorithm for the columns in each micro-partition.

Important

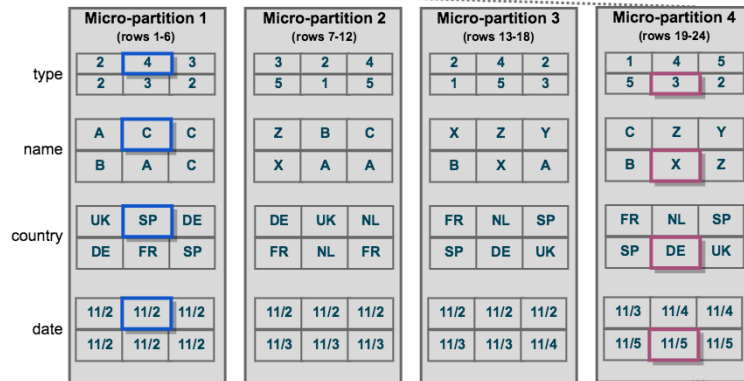
In Snowflake, as data is inserted/loaded into a table, clustering metadata is collected and recorded for each micro-partition created during the process. Snowflake then leverages this clustering information to avoid unnecessary scanning of micro-partitions during querying, significantly accelerating the performance of queries that reference these columns.

Table: t1

Logical Structure

type	name	country	date
2	A	UK	11/2
4	C	SP	11/2
3	C	DE	11/2
2	B	DE	11/2
3	A	FR	11/2
2	C	SP	11/2
3	Z	DE	11/2
2	B	UK	11/2
4	C	NL	11/2
5	X	FR	11/3
1	A	NL	11/3
5	A	FR	11/3
2	X	FR	11/2
4	Z	NL	11/2
2	Y	SP	11/2
1	B	SP	11/3
5	X	DE	11/3
3	A	UK	11/4
1	C	FR	11/3
4	Z	NL	11/4
5	Y	SP	11/4
5	B	SP	11/5
3	X	DE	11/5
2	Z	UK	11/5

Physical Structure



The table consists of 24 rows stored across 4 micro-partitions, with the rows divided equally between each micro-partition. Within each micro-partition, the data is sorted and stored by column, which enables Snowflake to perform the following actions for queries on the table:

1. First, prune micro-partitions that are not needed for the query.
2. Then, prune by column within the remaining micro-partitions.

Note that this diagram is intended only as a small-scale conceptual representation of the data clustering that Snowflake utilizes in micro-partitions. A typical Snowflake table may consist of thousands, even millions, of micro-partitions.